Revised memory on new held-out paths

Both policies compared with the same validation-selected fixed baseline; lower loss is better

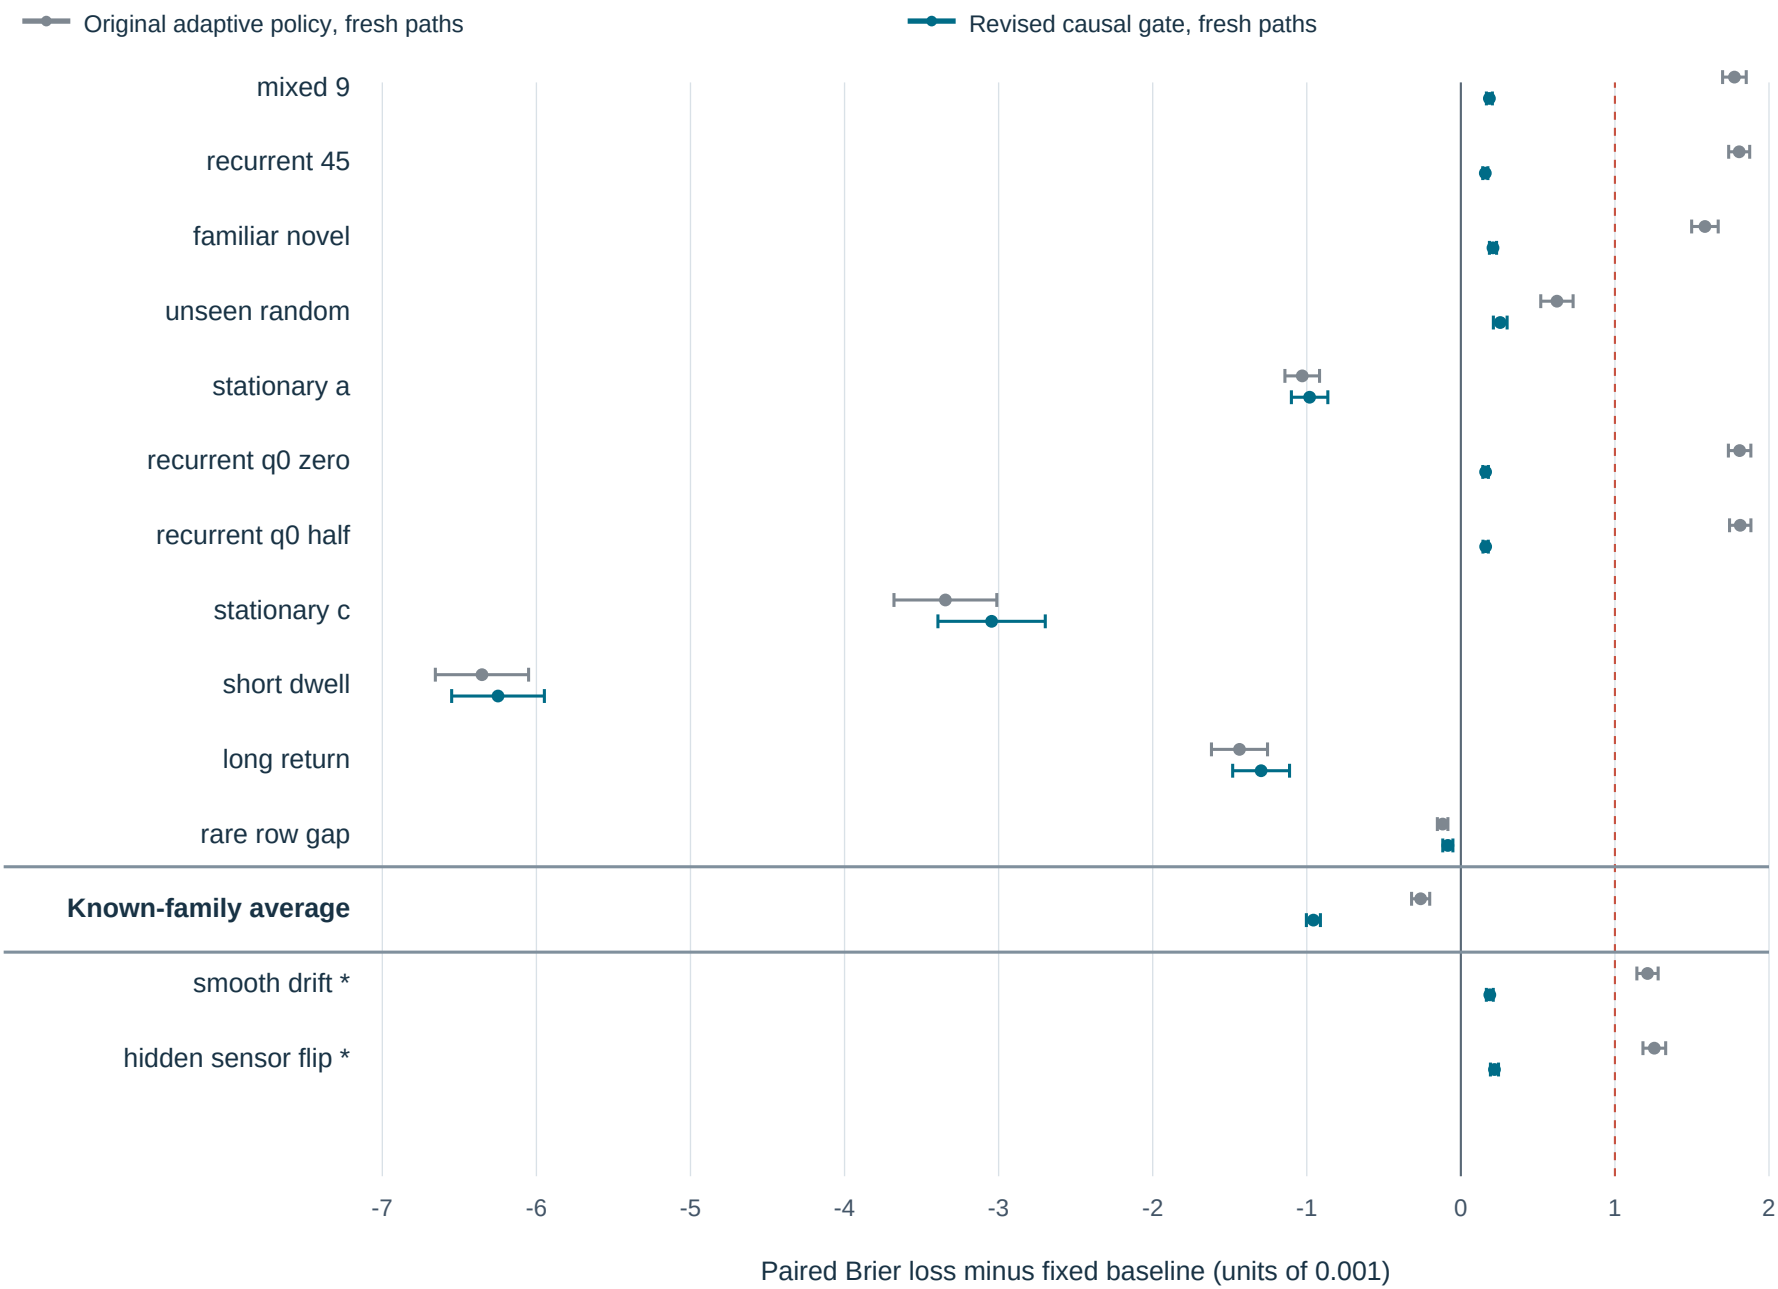

Bars: pointwise 95% Monte Carlo intervals. Red line: +0.001 harm margin, formally tested on known families only.

* New exploratory families: excluded from tuning and both primary gates; the fixed policy remains better on these families.